



Data Science Workflow – Flight Fare Prediction



Author: Anwesha Das

Created with Pi

CONTENTS

1. Understanding the Problem

3. Preprocessing Data

5. Exploratory Data Analysis (EDA)

7. Deployment

2. Collecting Data

Feature engineering means creating new, more meaningful inputs for the model. Some useful approaches include:

6. Model Building



1. Understanding the Problem

Goal: Predict the ticket price for a flight.

Why: Helps airlines set prices and customers plan better.

Question: “Given details of a flight (airline, source, destination, date, time, etc.), what will be its fare?”

2. Collecting Data

Ways to get data:

Airline websites

scrape prices.

Travel portals (e.g., MakeMyTrip, Skyscanner)

collect flight info.

Public datasets

Kaggle or government portals.

Company's own booking records

if working with airlines.

Collected data usually has:

Airline name

Date of journey

Departure and arrival time

Duration of flight

Source & destination

Number of stops

Ticket price



3. Preprocessing Data

1

Remove missing values (e.g., if price is not available).

2

Handle wrong entries (e.g., negative duration).

3

Standardize date and time format.

Feature engineering means creating new, more meaningful inputs for the model. Some useful approaches include:

Date & Time Features

- Extract day, month, weekday, and weekend/holiday flags.
- Convert departure/arrival times into categories (morning, afternoon, evening, night).

Duration Features

- Convert total flight duration into minutes.

Categorical Encoding

- Encode airline names, source, and destination into numerical form (label encoding or one-hot encoding).
- Represent number of stops as integers (non-stop = 0, one-stop = 1, etc.).

Scaling & Transformation

- Normalize or standardize features like duration and price.
- Apply log transformation to reduce skewness in ticket prices.



5. Exploratory Data Analysis (EDA)

Find patterns and trends in the data. Examples:

1

Prices may be higher for certain airlines.

2

Evening flights may cost more than morning flights.

3

Non-stop flights usually cost more than 1-stop flights.

4

Fares may increase during holiday seasons.

Use graphs and charts to understand these trends.



6. Model Building

Split data into train set (to learn) and test set (to check accuracy).

Choose models:

- Linear Regression: simple and fast.
- Random Forest / XGBoost: handles complex relations well.

Train the model with features → predict fares.

Evaluate performance using MAE, RMSE, R^2 (error metrics).



7. Deployment

1

Put the trained model into a system (website or app).

2

Example: A user enters details → Model predicts fare instantly.

3

Use tools like Flask, Django, or FastAPI to connect model with website.

4

Update model regularly with new data (as prices keep changing).

The background of the slide is a light blue color with a subtle, abstract pattern. It consists of numerous thin, wavy white lines that flow across the frame, interspersed with small, faint white dots, creating a sense of movement and depth.

Thank You